

Visualizing the distribution of a dataset

- ▶ When dealing with a set of data, often the first thing you'll want to do is get a sense for how the variables are distributed.
- ▶ This chapter of the tutorial will give a brief introduction to some of the tools in seaborn for examining univariate and bivariate distributions.
- ▶ You may also want to look at the categorical plots chapter for examples of functions that make it easy to compare the distribution of a variable across levels of other variables.

```
%matplotlib inline
import numpy as np
import pandas as pd
from scipy import stats, integrate
import matplotlib.pyplot as plt
import seaborn as sns
sns.set(color_codes=True)
np.random.seed(sum(map(ord, "distributions")))
```

The most convenient way to take a quick look at a univariate distribution in seaborn is the `distplot()` function. By default, this will draw a histogram and fit a kernel density estimate (KDE).

```
x = np.random.normal(size=100)
sns.distplot(x);
```





Figure:

Histograms

- ▶ Histograms are likely familiar, and a hist function already exists in matplotlib.
- ▶ A histogram represents the distribution of data by forming bins along the range of the data and then drawing bars to show the number of observations that fall in each bin.

- ▶ To illustrate this, let's remove the density curve and add a rug plot, which draws a small vertical tick at each observation.
- ▶ You can make the rug plot itself with the `rugplot()` function, but it is also available in `distplot()`:

```
sns.distplot(x, kde=False, rug=True);
```



When drawing histograms, the main choice you have is the number of bins to use and where to place them. `distplot()` uses a simple rule to make a good guess for what the right number is by default, but trying more or fewer bins might reveal other features in the data:

```
sns.distplot(x, bins=20, kde=False, rug=True);
```

